

Real-world model grades

Receiving, Prep, Pack, Returns, and Recovery. Each figure is that agent's own photo or dock set. They are not one combined accuracy.

Receiving · 12 dock deliveries

Prep · 15 products

Pack · 12 + 14

Returns · 40 scored

Recovery · no match rate

What you can put on a slide

Four agents have a model grade on real photos or real dock deliveries. Recovery has rules on six copied fee lines and no labeled match rate.

AGENT	SET	LABELERS	CALLS	HEADLINE
Receiving	12 real dock deliveries	Jaanu + Soanli	12	Full agreement on all 6 checks: 3 units. Quantity 10/10 where both agreed.
Prep	15 products	One frozen sheet	15	104/122 fields (85.2%). FNSKU value 15/15.
Pack · Gemini	12 warehouse carton photos	Expected manifests	12	5/12 decisions (41.7%). 0 false seals.
Pack · Claude	14 open-box orders	One sheet	15*	3/14 verdicts (21.4%). 0 false seals.
Returns	40 agreed, graded cases	Rishik Goud + Lasya	47	Identity 92.5% · Completeness 75.0% · Condition 42.5%
Recovery	6 copied fee lines	None	0	No match rate. No line written as claimable USD.

*Pack Claude called 15 orders, then removed ORD-005 from the score. Sources: pod eval 7 Oct, Prep 8 Oct, Pack Gemini 8 Oct, Pack Claude 9 Oct, Returns freeze 1 Oct.

40 agreed cases with a model grade

50 cases. Two labelers before scoring. Agreed means identity, completeness, and Amazon condition all match. Three agreed cases had no photos, so the rates use 40. One model call per graded return. policy_v1 chooses disposition.

37/40

Identity · 92.5%

0 false positives · 0 false negatives

30/40

Completeness · 75.0%

2 false positives · 0 false negatives

17/40

Condition · 42.5%

Amazon grade equality

0

Model errors

47 calls · about \$0.72

CHECK	MATCH	UNCERTAIN	FALSE POSITIVES	FALSE NEGATIVES
Identity	37/40 (92.5%)	1/40 (2.5%)	0	0
Completeness	30/40 (75.0%)	5/40 (12.5%)	RTN-046, RTN-049	0
Condition	17/40 (42.5%)	Not a pass/fail count	Not a pass/fail count	Not a pass/fail count

Seven disagreements stay out of every rate: RTN-002, 005, 006, 009, 013, 037, 044. No-photo cases: RTN-017, 026, 032. Cohen’s kappa was not computed.

The usual miss is one Amazon step

CONDITION MISMATCH ON THE 40	COUNT
Human Very Good, agent Like New	14
Human Good, agent Very Good	4
Human Like New, agent Very Good	4
Human unknown, agent Very Good (RTN-049)	1

23 condition mismatches. Identity non-matches that are neither false positives nor false negatives: RTN-028 (agent uncertain), RTN-010 and RTN-049 (labelers uncertain).

SAVED DISPOSITION ON 47 GRADES	COUNT
restock	38
pending_review	8
refurbish	1
liquidate	0
dispose	0

policy_v1 reproduced 40/40 saved dispositions from the agent's own checks. Against the agreed human checks, 27/40 (67.5%) match. Those 13 gaps are upstream check disagreements. Disposition was not human-labeled.

12 real dock deliveries

One Claude observation each, then receiving_policy_v1. Jaanu and Soanli labeled independently. Both agreed on every check for 3 units: UNIT-RCV-01, 06, 10. Primary rates use those 3. All 12 outcomes were pending_review.

3/12

Full agreement on all six checks

10/10

Quantity, where both labelers agreed

11/12

Unit damage, soft agreement

0

Ingestion failures

PRIMARY · N = 3

CHECK	MATCH	FP	FN
Identity	1/3	0	0
Carton count	0/3	0	0
Quantity	3/3	0	0
Carton damage	2/3	0	1
Unit damage	3/3	0	0
Quality flags	3/3	0	0

SOFT · HUMANS AGREE ON THAT CHECK

CHECK	N	MATCH
Identity	12	5/12 (41.7%)
Carton count	6	3/6 (50%)
Quantity	10	10/10
Carton damage	8	6/8 (75%)
Unit damage	12	11/12 (91.7%)
Quality flags	11	7/11 (63.6%)

Clear paths: 07 and 08 wrong product. 09 quantity incomplete. 04, 05, 11 carton count below PO. 12 too blurry. Soft rates do not replace the primary 3. This is a 12-delivery pilot.

eval/EVALUATION.md · scored 2026-10-07T17:34Z

05

104 of 122 fields agree

claude-sonnet-4-5 on PRODUCT-1 through PRODUCT-15. One frozen sheet. 201.6 seconds. Uncertain is not counted as a match. Polybag seal and suffocation warning are not scorable on this sheet.

85.2%

104 agreements / 122 scorable fields

15/15

FNSKU visible and FNSKU value

15/15

Polybag present

7

Disagreements · 11 model-uncertain

FIELD	SCORABLE	AGREE	DISAGREE	UNCERTAIN
Barcode visible	15	10	0	5
Barcode readable	14	7	2	5
FNSKU visible / value	15	15	0	0
Polybag present	15	15	0	0
Expiry visible	15	14	1	0
Expiry readable	3	2	0	1
Handling mark	15	14	1	0
Photo quality	15	12	3	0

Disagreements: PRODUCT-3 expiry, PRODUCT-5 and 12 barcode readable plus photo quality, PRODUCT-7 photo quality, PRODUCT-15 handling mark. A later note on PRODUCT-15 and PRODUCT-3 was not used to rescore.

5 of 12 decisions, zero false seals

Twelve warehouse carton photos. Order-blind Gemini perception, then Python stages 2-5. When the pipeline is unsure, it returns UNCERTAIN. It did not seal a bad carton.

5/12

Correct decisions · 41.7%

0

False seals

2/2

SEAL precision

1/1

STOP_FIX precision

EXPECTED	PREDICTED SEAL	PREDICTED STOP_FIX	PREDICTED UNCERTAIN	TOTAL
SEAL	2	0	3	5
STOP_FIX	0	1	4	5
UNCERTAIN	0	0	2	2

Matches: correct apparel, blurry hold, unknown product hold, correct electronics, wrong electronics caught. The seven misses are catalogue ambiguity (score 0.80 under a 0.85 threshold). Stage 1 perception failures: 0. No tuning after this baseline.

3 of 14 verdicts, still zero false seals

claude-sonnet-4-5, 23 photos, about 200 seconds. ORD-005 was removed because the photo includes a person’s head. One labeler. The agent sealed nothing.

3/14

Verdict match · 21.4%

0

False seals

13/13

Product damage agreement

10/13

Missing-item agreement

PATTERN	ORDERS	WHAT HAPPENED
Label FAIL, agent FAIL	3	ORD-002, ORD-010, ORD-021. Extra items. These are the three matches.
Label PASS, agent FAIL	5	ORD-004, 027, 031, 042, 050. Extras, or missing plus extra.
Label PASS, agent UNCERTAIN	2	ORD-015, ORD-022. Low confidence or unclear brand.
Label FAIL, agent UNCERTAIN	3	ORD-018, 028, 044.
Label UNCERTAIN, agent FAIL	1	ORD-040. Variant or colour mismatch.

Extra-item agreement is 3/14. Cables, remotes, manuals, and stands that belong with the product were treated as extras. This Claude sheet and the Gemini 12 are separate grades.

Rules ran. There is no accuracy number.

The decision is the rule code. A model summary does not change the assessment. Six copied fee lines were judged locally with zero model calls. No line was written as claimable USD.

0

Model calls on these lines

6

Fee lines judged

2

SUPPORTS · PRODUCT-1, PRODUCT-8

4

SILENT · PRODUCT-4, 6, RTN-019, iPhone

CHARGE TYPE	WHAT THE RULE DOES
Inbound defect fee	Prep PASS contradicts the fee. Prep FAIL supports it. Prep UNCERTAIN stays silent.
Refund, item not returned	Completed Returns identity PASS contradicts. Identity FAIL supports. Pending stays silent.
Weight tier / lost inbound	Silent. No demo measurement. A receiving shortfall is not a channel-loss claim.
Amount missing or 0.00	A CONTRADICTS result is forced back to silent.

Order OD435816013145832100 stays UNCERTAIN: the handset line is missing org, unit, SKU, return event, recipient, and due date. A match rate needs human labels of CONTRADICTS, SUPPORTS, or SILENT on real fee lines. That sheet does not exist.

CLOSE

Three lines to leave on screen

92.5%

Returns identity on 40 agreed cases

85.2%

Prep fields on 15 products

0

False seals on both Pack runs

SAY THIS	LEAVE THIS OFF THE SLIDE
Returns identity is strong on this set. Condition sits one Amazon step high.	An overall Pod 4 accuracy, or a claim that the same rates hold next week.
Pack fails closed. Both real-photo runs recorded zero false seals.	Adding 41.7% and 21.4% into one Pack score.
Receiving quantity matched on every row both labelers agreed.	Treating the soft table as the primary 50-case evaluation.
Recovery withholds a claim when the fee line is incomplete.	A recovery precision or a dollar saved.

Frozen files: Returns eval/frozen/2026-10-01T1015Z. Receiving eval/frozen/2026-10-07T1734Z. Prep held-out Claude results. Pack REAL_WORLD_EVALUATION.md and held-out Claude results.