

SYDON.AI

× CODEQUESTERS



Cube Buildathon

OFFICIAL PARTICIPANT HANDBOOK

Round 2 — Individual Build • Rules • Submission • Evaluation

DON'T JUST BUILD AI. ENGINEER IT.

Welcome to Cube Buildathon

CUBE is a production-oriented AI engineering challenge by **Sydon.AI** in collaboration with **CodeQuesters**. You are now entering the most important part of the experience: **Round 2 — Build**.

This handbook is your operating guide for the individual build phase. Read it before you start, keep it open while building, and use it to verify your submission before the deadline.

CUBE = Collaborate · Uplift · Build · Earn

The core idea

Build one focused AI agent around your selected real-world commerce problem. Engineer it so that another person can understand the workflow, reproduce the result, inspect the evidence and trust the boundaries of the system.

AI tools can accelerate your work. Engineering judgment stays yours.

1. Round 2 Starts Now

25 September 2026 • 9:00 AM IST marks the official start of the Round 2 build phase. Your problem statement is already selected. From this point, the goal is execution.

What you do next

- Re-read your selected problem statement completely.
- Identify inputs, required checks, expected outputs, constraints and edge cases.
- Open the official repository for your track and fork it into your own GitHub account.
- Set up the project locally and prove that the basic environment works.
- Read the repository README, RULES and getting-started instructions.
- Design the workflow: input → processing → checks → decision → evidence → output.
- Get a small end-to-end example working early, then improve it.
- Start evaluating while you build instead of waiting until the final day.

Round 2 is individual

You are building one specialised agent independently. Your Round 2 score is out of 100 and will matter later if you progress to Round 3.

2. The Five Problem Statements

Each participant selects exactly one track. Your selected track defines the scope of your individual build.

TRACK	AGENT	ROUND 2 FOCUS
01	Receiving Manager	Reason over receiving/inbound evidence, identify operational discrepancies or issues, and produce an evidence-backed decision.
02	Prep Manager	Evaluate whether products meet the defined preparation requirements before the next operational stage.
03	Pack Manager	Assess packaging readiness and determine whether an item can proceed, needs correction, or should stop.
04	Returns Manager	Assess returned-item condition, classify the outcome and create structured evidence for downstream decisions.
05	Recovery Manager	Read upstream evidence and reason about recoverable charges/claims using evidence rather than redoing upstream image classification.

Do not dilute your build by trying to solve all five tracks. Build your selected agent deeply and make its boundaries clear.

3. Hard Timeline & Submission Deadline

WHEN	STAGE	MEANING
25 Sep • 9:00 AM IST	BUILD OPENS	Round 2 individual build officially begins.
27 Sep	SUBMISSIONS OPEN	The official submission form begins accepting final submissions.
1 Oct • 6:00 PM IST	FINAL DEADLINE	The submission form closes permanently. No reopening and no extension.
2–3 Oct	TECHNICAL EVALUATION	Round 2 submissions are evaluated using the five-criterion rubric.
4 Oct	ROUND 3 FINALISTS	Selected participants are announced.
5–9 Oct	ROUND 3 • INTEGRATE	Selected finalists work in five-person Pods.
10 Oct	GRAND FINALE	Final teams present and defend their integrated system.

■ **FINAL MEANS FINAL:** Everything required for Round 2 must be complete and submitted by 1 October 2026, 6:00 PM IST. Once the form closes, it will not be reopened.

4. GitHub & Build-Phase Rules

Your repository history is part of the engineering context. Build cleanly and keep the final codebase reproducible.

Required workflow

- Fork the official track repository into your own GitHub account.
- Clone your fork locally and configure the environment.
- Build, test, document and deploy from your fork.
- Push your work to your fork during the official build phase.
- Use meaningful commits so your progress is understandable.

Commit cut-off

- **All Round 2 code commits must be made during the official build phase.**
- After the build phase ends, do not continue committing code to the Round 2 repository.
- Do not make post-build code changes and submit an earlier state.
- If you discover an issue later, document it rather than silently altering the submitted code.

Important distinction

The **submission deadline** is 1 October at 6:00 PM IST. The organisers will communicate the exact operational cut-off for build activity if it differs from the submission window. In all cases, your submitted code must reflect work completed within the authorised build phase.

5. What You Must Submit

Submissions open from **27 September**. Treat the form as the final packaging step for work completed during Round 2.

- **GitHub repository:** your fork with the complete Round 2 implementation.
- **README.md:** problem understanding, solution overview, setup, usage, assumptions and limitations.
- **ARCHITECTURE.md:** architecture, components, data flow, model/agent usage and important engineering decisions.
- **Demo video:** a concise demonstration of the working system and important decisions.
- **Deployment URL:** where applicable.
- **Evaluation report/results:** methodology, measurements, failure modes and uncertainty handling.
- **LinkedIn post URL:** mandatory for Round 2.
- **Any additional links or files explicitly requested in the official submission form.**

No resubmission

Once you submit, your submission is final. There will be no resubmission facility. Check every field and link before clicking submit.

6. Mandatory LinkedIn Post

Every Round 2 participant must publish a LinkedIn post about their CUBE build and include the post URL in the submission form.

- State your selected CUBE track.
- Explain what your agent does and the operational problem it addresses.
- Share a meaningful engineering detail, result, demo or learning.
- Tag **CodeQuesters** and **Sydon.AI**.
- Use the official CUBE context/hashtags communicated by the organisers.
- Paste the live LinkedIn URL into the submission form.

Official template

The organisers will share a LinkedIn post template shortly, targeted for **27–28 September**. You can use it as the base, but your technical content should reflect your own build.

7. Round 2 Evaluation — 100 Points

Five criteria. One hundred points. The rubric rewards problem understanding, working functionality, measured performance, traceability and clear delivery.

NO.	CRITERION	PTS	WHAT IS ASSESSED
1	Problem Understanding & Solution Relevance	15	Understanding the assigned problem, operational relevance, scope and assumptions.
2	Agent Functionality & Decision Quality	25	Working agent, required checks, input handling, processing, structured outputs, edge cases and decisions.
3	Evaluation, Accuracy & Uncertainty Handling	25	Evaluation methodology, measured performance, false positives/negatives, UNCERTAIN handling and failure modes.
4	Evidence, Traceability & Engineering Quality	20	Evidence, per-check verdicts, confidence, model/version, timestamps, overrides, architecture, reliability and security.
5	UX, Demo & Documentation	15	Usability, workflow clarity, decision/evidence visibility, demo, README, architecture docs, deployment and required links.

TOTAL — 100 POINTS

8. What Evaluators Look For

Problem Understanding & Solution Relevance — 15 points

- Explain the actual problem clearly.
- Map the solution to the operational workflow.
- Keep assumptions explicit.
- Avoid irrelevant features.

Agent Functionality & Decision Quality — 25 points

- Implement required track checks.
- Handle inputs and edge cases.
- Return structured outputs.
- Make decisions traceable to checks/evidence.

Evaluation, Accuracy & Uncertainty Handling — 25 points

- Measure rather than claim.
- Use held-out data where applicable.
- Report false positives/negatives and uncertainty.
- Document failure modes honestly.

Evidence, Traceability & Engineering Quality — 20 points

- Capture evidence behind important decisions.
- Include per-check verdicts and confidence.
- Record overrides/reasons where applicable.
- Show clean architecture, error handling and maintainability.

UX, Demo & Documentation — 15 points

- Make the workflow easy to follow.
- Show input → decision → evidence.
- Make setup reproducible.
- Use the demo to show the real working system.

9. Evidence & Decision Traceability

CUBE evaluates engineering systems, not screenshots. Your agent should make it possible to understand what it saw, what it checked, what it decided and why.

The fixed evidence contract includes fields such as:

- record_id and schema_version
- organization_id and client_id
- agent and subject
- captured_at and operator_label
- images where applicable
- checks[] with check_key, verdict, confidence, detail, model_version and latency_ms
- outcome with decision, decided_by and decided_at
- overrides[]
- status and content_hash

PASS / FAIL / UNCERTAIN

- **PASS:** the available evidence supports the condition.
- **FAIL:** the available evidence supports that the condition is not met.
- **UNCERTAIN:** the evidence does not support a reliable judgment.
- UNCERTAIN is not a low-confidence PASS. Use it when evidence is insufficient or genuinely ambiguous.

Recovery Manager

Recovery should reason over upstream evidence records and report claim/charge precision and evidence correctness rather than treating the task as a fresh image-classification benchmark.

10. Evaluation Methodology

Evaluation quality is part of your score. A strong submission demonstrates how you know whether the agent works.

- Use an unseen/held-out evaluation set where applicable.
- For vision tracks, use at least **50 unseen units** where applicable.
- Have two human evaluators independently label units before the agent runs.
- Report human agreement where possible, such as Cohen's kappa.
- Include varied conditions: good/poor lighting, different angles, slight blur and genuinely ambiguous cases.
- Report per-check performance where applicable.
- Report false positives, false negatives, UNCERTAIN rate, latency/cost where relevant and failure modes.
- Prefer honest measured results and documented limitations over unsupported accuracy claims.

A useful evaluation table

Make it easy to see: test unit → human label → agent result → agreement/disagreement → uncertainty → failure-mode notes.

11. Round 2 → Round 3

Round 2 is about **individual performance**. Round 3 is about **team integration**.

Round 2 — Individual Build

- One participant builds one specialised agent.
- Round 2 is scored out of 100.
- Your individual implementation, evaluation, evidence, UX and documentation are assessed.

Round 3 — Pod Integration

- Selected participants form **five-person Pods**.
- A Pod brings together Receiving, Prep, Pack, Returns and Recovery.
- The team integrates the five independently built agents into one connected end-to-end commerce system.
- Round 3 focuses on collaboration, integration, interfaces/evidence, dependencies and system-level execution.

Your Round 2 score carries forward

For participants who reach Round 3, Round 2 and Round 3 scores are combined for the final result. Your individual performance now therefore contributes to your final standing later.

12. Final Scoring & Prize Pool

Final scoring combines the two build stages for participants who reach Round 3.

STAGE	MODE	MAX	FOCUS
ROUND 2	INDIVIDUAL BUILD	100	Individual agent and engineering performance.
ROUND 3	POD INTEGRATION	100	Team integration and system-level execution.
FINAL	COMBINED	200	Round 2 + Round 3.

■50,000 TOTAL CASH PRIZE POOL

WINNER	■25,000
1ST RUNNER-UP	■15,000
2ND RUNNER-UP	■10,000

The total prize pool is ■50,000: ■25,000 for the Winner, ■15,000 for the 1st Runner-Up and ■10,000 for the 2nd Runner-Up.

13. Final Submission Checklist

Before clicking Submit, verify every line.

- Correct selected track.
- Correct GitHub fork and final code.
- Required code commits were made during the authorised build phase.
- README.md complete.
- ARCHITECTURE.md complete.
- Demo video works and clearly shows the solution.
- Deployment URL works, if applicable.
- Evaluation methodology/results documented.
- Evidence and decision trace are understandable.
- LinkedIn post published.
- CodeQuesters and Sydon.AI tagged.
- LinkedIn post URL copied correctly.
- All submission-form fields completed.
- No missing links/files.
- You understand that the submission cannot be resubmitted.
- You will submit before **1 October 2026, 6:00 PM IST**.

Build → Test → Measure → Document → Publish → Submit.

14. Common Mistakes to Avoid

- Building before understanding the operational workflow.
- Making a generic chatbot instead of a focused agent.
- Forcing ambiguous cases into PASS or FAIL instead of using UNCERTAIN.
- Showing UI without proving decision quality.
- Skipping evaluation or using only cherry-picked examples.
- Making code changes after the authorised build-phase cut-off.
- Broken setup instructions or incomplete documentation.
- Forgetting the mandatory LinkedIn post.
- Assuming you can replace a submission later.
- Trying to solve all five tracks instead of doing one deeply.

15. FAQ

Can I change my selected problem statement?

No. Your selected track defines your Round 2 scope.

Is Round 2 individual?

Yes. Round 2 is an individual build.

When does the build phase start?

25 September 2026 at 9:00 AM IST.

When do submissions open?

27 September 2026.

What is the final deadline?

1 October 2026 at 6:00 PM IST.

Will the form reopen?

No. The deadline is final.

Can I resubmit?

No. There is no resubmission facility.

Can I commit code after the build phase?

No. Round 2 code commits must be made during the authorised build phase.

Is LinkedIn mandatory?

Yes. A LinkedIn post and its URL are mandatory for Round 2.

Who must I tag?

CodeQuesters and Sydon.AI.

How is Round 2 scored?

Five criteria, total 100 points.

Does Round 2 affect the final result?

Yes. For Round 3 participants, Round 2 and Round 3 scores are combined.

What happens in Round 3?

Selected participants form five-person Pods and integrate the five specialised agents into one connected system.

16. Final Word to Builders

The goal of Round 2 is not the biggest codebase. It is a focused, working, measurable and defensible engineering system.

You already chose the problem. Now make the next decisions count: understand the workflow, design the right agent, build deliberately, test honestly, capture evidence and document what you learned.

Round 3 will ask you to integrate your work with four other builders. A strong individual system with clear interfaces, evidence and engineering decisions is your foundation for that stage.

CUBE moves from individual ownership to system-level collaboration. Your Round 2 build is the foundation.

Cube Buildathon
SYDON.AI × CODEQUESTERS

DON'T JUST BUILD AI. ENGINEER IT.